

Idea 1 introduction contribution

Our first analysis examines the premise behind progressive cross-modal norm supervision: whether perceptual observations improve normative decisions and whether the improvement survives answer-text predictability. Using released EgoNormia predictions and direct audits of the proposed training resources, we compare matched input conditions, inspect norm-specific differences, and test linguistic shortcuts. We also examine contextual language and visual representations. These analyses refine the curriculum’s controls and data preparation; they do not yet establish that sequential training transfers better than an equal-composition mixture.

Idea 1 discussion contribution

Implications for Idea 1. The evidence supports retaining a grounded supervision stage, while narrowing what the current analysis establishes. Released frame-grid predictions outperform description predictions on matched items, including items missed by a text-only probe. However, a potential description-prompt defect prevents interpreting the joint-score gap as a clean measure of perceptual necessity. Before training, we will verify the target backbone’s frame handling and rerun matched, consistently prompted controls. We will preserve the original comparison of the curriculum with an equal-composition unordered mixture and order ablations; the present analysis does not test that central hypothesis.

Three data decisions follow. First, deduplicate NormBank and explicitly resolve conflicting context labels before splitting or sampling. Second, preserve NormLens’s native judgment task and annotator disagreement rather than generating stylistic multiple-choice distractors. Third, use probe-error results alongside full-set results, with text-only controls and equal temporal evidence for video versus frames. A supervised EgoNormia diagnostic probe is not the transfer model, but it does use benchmark labels: we should describe the benchmark as excluded from transfer-model training, not as untouched by analysis. The team has contacted the VideoNorms authors to request annotations and expects access soon. Annotation analysis remains pending receipt; no result here establishes its training value. The pooled CLIP diagnostic and selected cases further motivate representations that preserve interaction details and distinguish action choice from justification quality.

AI usage and contribution disclosure

Scope of this disclosure. This statement covers the supplied Analysis 1 package. It must be incorporated into the team’s complete disclosure; it does not certify other teammates’ work.

The prior package and recommendations were AI-assisted, as recorded in the supplied conversation excerpts. Exact earlier model versions and complete usage history were not provided. For this revision, OpenAI Codex was used substantially: it inspected the source files and assignment, retrieved public data and pretrained weights, wrote and ran analysis and validation code, identified statistical and implementation caveats, generated figures and LaTeX sections, inspected qualitative images, and conducted a strict claim audit. The draft is therefore not appropriately described as AI assistance limited to wording. Human authors must verify the scientific interpretation and take responsibility for the submitted work.

Additional pretrained models used computationally were sentence-transformers/all-mpnet-base-v2 for frozen contextual text features and openai/clip-vit-base-patch32 for frozen image features. Their pinned revisions and feature-processing details are included in the artifact logs. Empath, TF-IDF, SVD, and fitted linear classifiers were used as diagnostic representations or probes. The VLM grounding results are rescored author-released predictions, not new VLM inference. The user supplied 60 responses for a 20-item, three-stage exploratory human exercise in chat, administered in two sequential ten-item blocks with feedback between blocks. Codex administered the prompts, transcribed the responses, scored action agreement against released gold, and assigned post-hoc evidence categories as a single AI coder. No independent second coder or second human participant was included. Unaided completion was not independently verified. AI inspection of the separate two qualitative examples is not human annotation. No full multimodal model was trained.

Human contributions. The supplied proposal identifies Siddharth Singh, Anisha Reddy, Atharv Patawar, Yash Lothe, and Shiv Gupta. The user supplied the human-exercise responses and reported contacting the VideoNorms authors for data access. The record does not establish each person’s complete Assignment 2 contributions. The team must add truthful individual descriptions of idea development, coding/debugging, literature review, and writing. No teammate contribution or TA consultation is fabricated in this package.