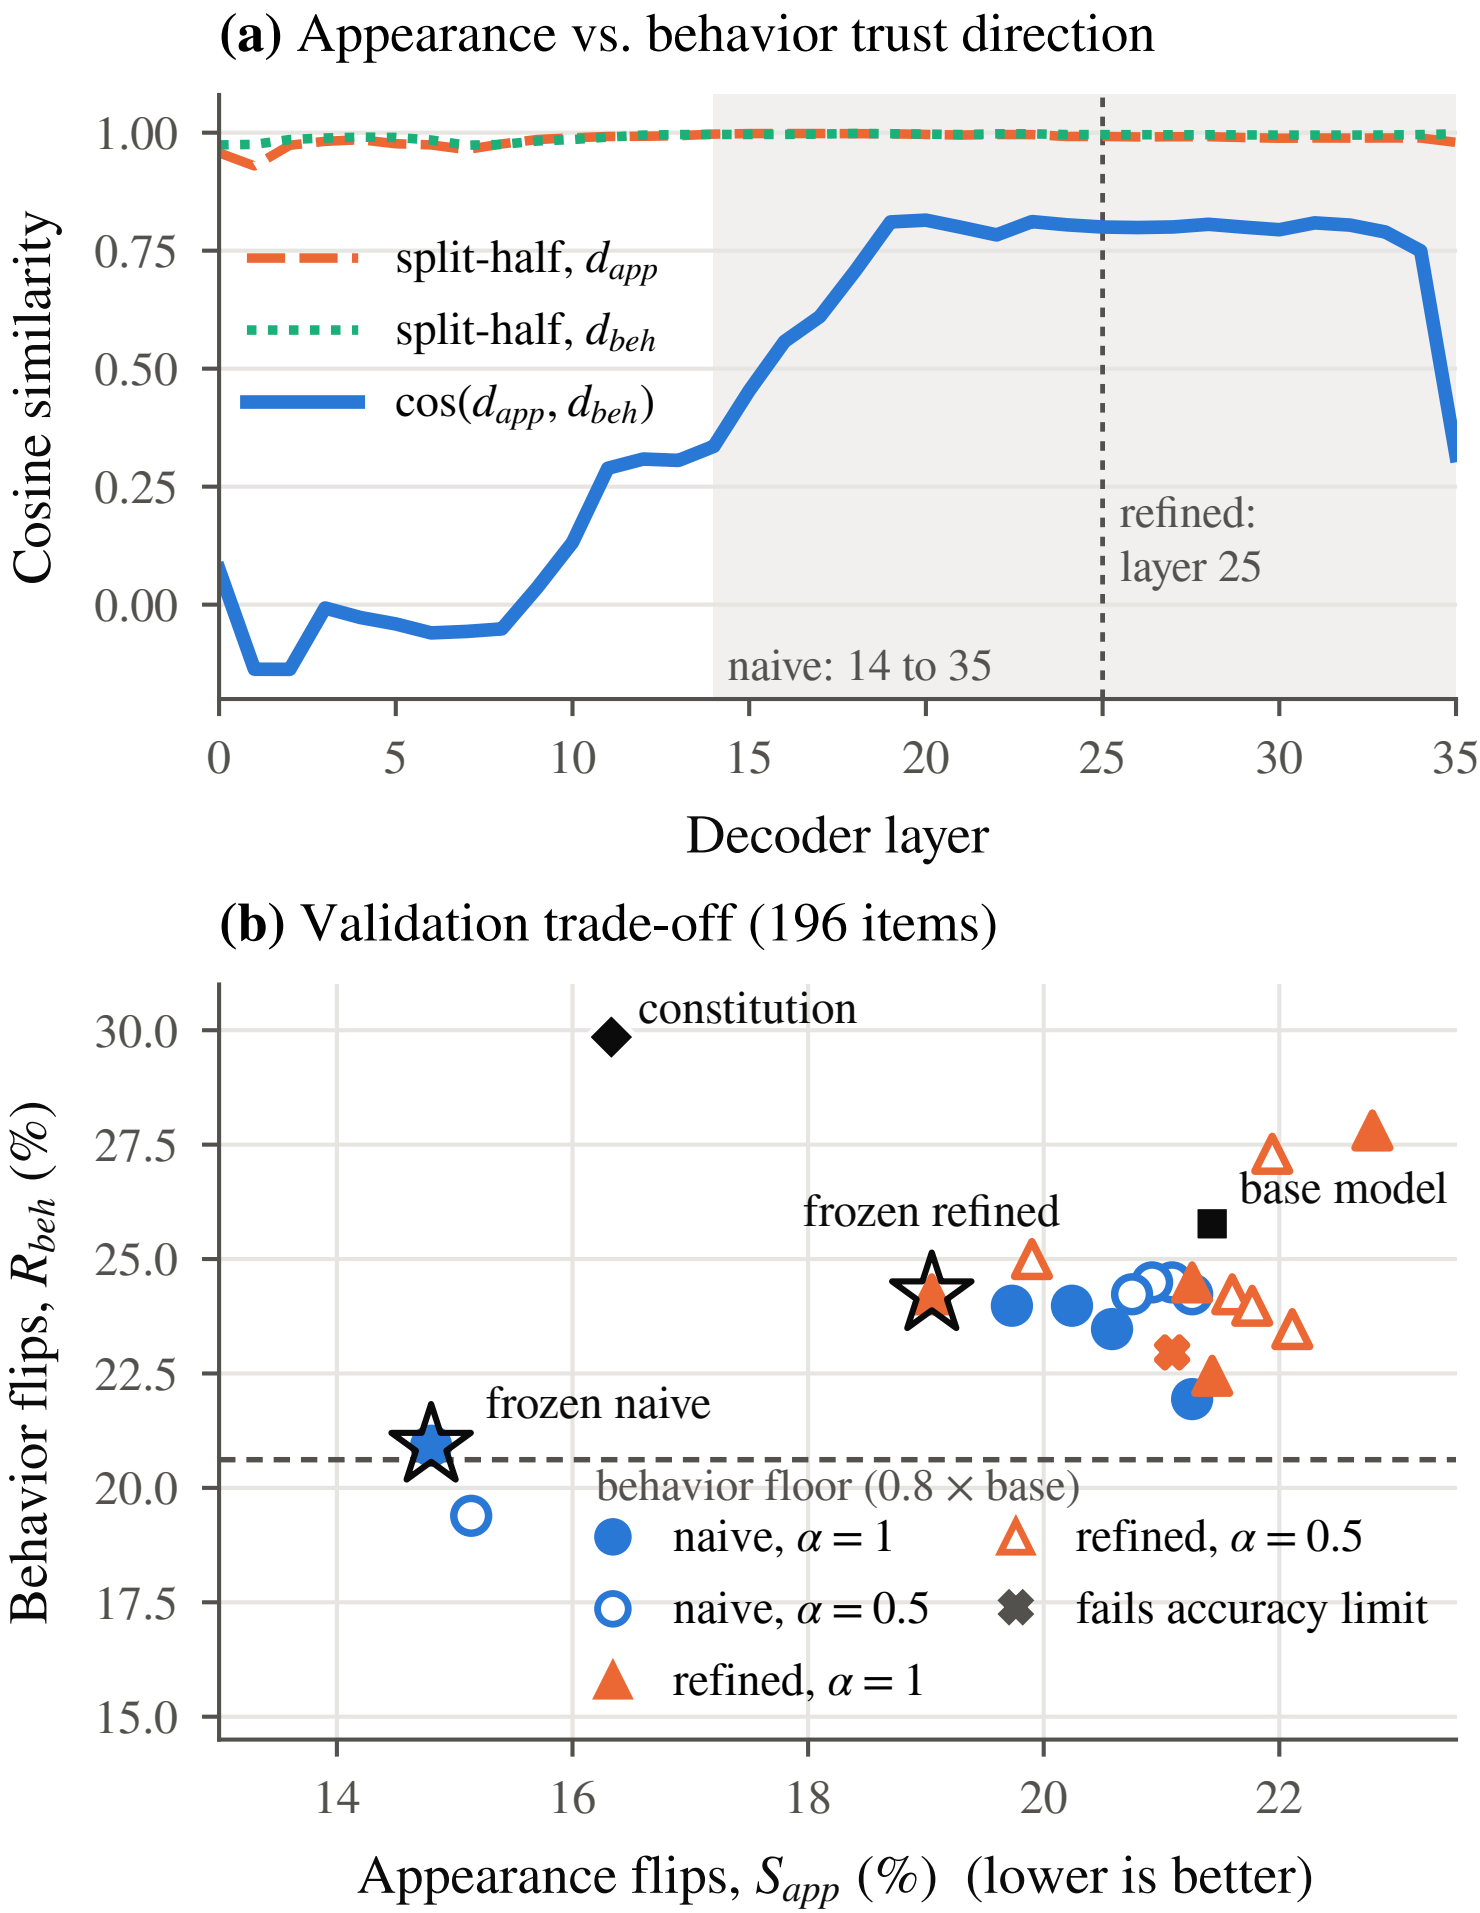
(a) Appearance vs. behavior trust direction
Cosine similarity
1.00
0.75
0.50
0.25
0.00
split-half, d_{app}
split-half, d_{beh}
cos(d_{app}, d_{beh})
naive: 14 to 35
refined:
layer 25
0
5
10
15
20
25
30
35
Decoder layer
(b) Validation trade-off (196 items)
Behavior flips, R_{beh} (%)
30.0
27.5
25.0
22.5
20.0
17.5
15.0
constitution
base model
frozen refined
frozen naive
behavior floor (0.8 × base)
naive, $\alpha = 1$
naive, $\alpha = 0.5$
refined, $\alpha = 1$
refined, $\alpha = 0.5$
fails accuracy limit
14
16
18
20
22
Appearance flips, S_{app} (%) (lower is better)